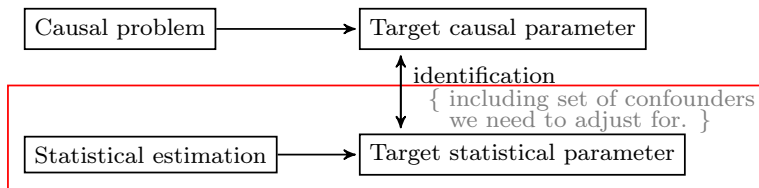


Estimating the target

Estimating the target



- ▶ one estimator is not more causal than another.
- ▶ different estimators are based on different nuisance parameters and have different statistical properties (bias/variance).

G-formula versus IP-weighting

- G-formula
1. Estimate nuisance parameters
 $f(a, x) = \mathbb{E}[Y \mid A = a, X = x]$ and the average over the marginal distribution μ_X of X
 2. Plug in to estimate the ATE:

$$\hat{\psi}_n^{\text{g-formula}} = \tilde{\Psi}(\hat{f}_n, \hat{\mu}_X) = \int_{\mathbb{R}^d} (\hat{f}_n(1, x) - \hat{f}_n(0, x)) d\hat{\mu}_X(x)$$

- IP-weighting
1. Estimate nuisance parameters
 $\pi(a \mid x) = P(A = a \mid X = x)$ and the average over the distribution P of O
 2. Plug in to estimate the ATE:

$$\hat{\psi}_n^{\text{ipw}} = \tilde{\Psi}_{\text{ipw}}(\hat{\pi}_n, \hat{P}_n) = \int_{\mathbb{R}^d} \left(\frac{ay}{\hat{\pi}_n(a \mid x)} - \frac{(1-a)y}{\hat{\pi}_n(a \mid x)} \right) d\hat{P}_n(x)$$

One-step estimation

One-step

1. Estimate nuisance parameters

$$f(a, x) = \mathbb{E}[Y \mid A = a, X = x], \pi(a \mid x) = \mathbb{E}[A \mid X = x]$$

and the average over the distribution P of O

2. Plug in to estimate the ATE:

$$\hat{\psi}_n^{\text{one}} = \tilde{\Psi}_{\text{one}}(\hat{f}_n, \hat{\pi}_n, \hat{P}_n) = \int_{\mathbb{R}^d} \sum_{a=0,1} \sum_{y=0,1} \left\{ \left(\frac{a}{\hat{\pi}_n(a \mid x)} - \frac{1-a}{\hat{\pi}_n(a \mid x)} \right) (y - \hat{f}_n(a, x)) + \hat{f}_n(1, x) - \hat{f}_n(0, x) \right\} d\hat{P}_n(o)$$

G-formula versus IP-weighting

Estimation of the averages over μ_X and P is straightforward using the empirical average over the observed data.

This yields:

G-formula estimator:
$$\hat{\psi}_n^{\text{g-formula}} = \frac{1}{n} \sum_{i=1}^n \{ \hat{f}_n(1, X_i) - \hat{f}_n(0, X_i) \}$$

IP-weighted estimator:
$$\hat{\psi}_n^{\text{ipw}} = \frac{1}{n} \sum_{i=1}^n \left\{ \frac{A_i Y_i}{\hat{\pi}_n(A_i | X_i)} - \frac{(1 - A_i) Y_i}{\hat{\pi}_n(A_i | X_i)} \right\}$$

One-step estimator:
$$\hat{\psi}_n^{\text{one}} = \frac{1}{n} \sum_{i=1}^n \left\{ \frac{A_i}{\hat{\pi}_n(A_i | X_i)} - \frac{(1 - A_i)}{\hat{\pi}_n(A_i | X_i)} (Y_i - \hat{f}_n(A_i, X_i)) + \hat{f}_n(1, X_i) - \hat{f}_n(0, X_i) \right\}.$$

G-formula versus IP-weighting versus one-step/TMLE

Properties of the different estimators:

G-formula versus IP-weighting versus one-step/TMLE

Properties of the different estimators:

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

G-formula versus IP-weighting versus one-step/TMLE

Properties of the different estimators:

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

One-step estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

G-formula versus IP-weighting versus one-step/TMLE

Properties of the different estimators:

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

One-step estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

- ▶ the one-step estimator and **the TMLE estimator** share the same large-sample properties.

G-formula versus IP-weighting versus one-step/TMLE

Properties of the different estimators:

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

One-step estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

- ▶ the one-step estimator and **the TMLE estimator** share the same large-sample properties.
- ▶ one thing they have in common is an additional bias correction leads to consistency if either \hat{f}_n or $\hat{\pi}_n$ is consistent (commonly known as **double robustness**).

G-formula versus IP-weighting versus TMLE

SMALL EXERCISE:

By the law of large numbers, the one-step estimator converges in probability to:

$$\mathbb{E}_{P_0} \left[\left(\frac{A}{\pi(A|X)} - \frac{1-A}{\pi(A|X)} \right) (Y - f(A, X)) + f(1, X) - f(0, X) \right] \quad (1)$$

where (f, π) denotes the limit of $(\hat{f}_n, \hat{\pi}_n)$. Compute the right hand side of (1) when

1. $f = f_0$ (i.e., the outcome regression is consistently estimated), and
2. $\pi = \pi_0$ (i.e., the propensity score is consistently estimated).

G-formula versus IP-weighting versus TMLE

TMLE estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

- ▶ consistent if either \hat{f}_n or $\hat{\pi}_n$ is consistent.

G-formula versus IP-weighting versus TMLE

TMLE estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

- ▶ consistent if either \hat{f}_n or $\hat{\pi}_n$ is consistent.

Some other things we will see/explore:

- ▶ if $\hat{\pi}_n$ is modeled correctly, the TMLE estimator will have lower variance than the IP-weighted estimator.

G-formula versus IP-weighting versus TMLE

TMLE estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

- ▶ consistent if either \hat{f}_n or $\hat{\pi}_n$ is consistent.

Some other things we will see/explore:

- ▶ if $\hat{\pi}_n$ is modeled correctly, the TMLE estimator will have lower variance than the IP-weighted estimator.
- ▶ the TMLE estimator may have larger variance than the g-formula estimator based on a correctly modeled \hat{f}_n (but it gives protection against the case that it is not).

G-formula versus IP-weighting versus TMLE

Can't we just construct a good g-formula estimator???

G-formula versus IP-weighting versus TMLE

Can't we just construct a good g-formula estimator???

- ▶ a logistic regression — great if correctly specified, but horrible if not.
- ▶ a random forest — properly tuned?

A random forest — properly tuned?

Predictive performance of an estimator can be measured in terms of some distance¹ between:

- 1) the observed outcome: Y_i
- 2) and the predicted conditional expectation: $\hat{f}_n(A_i, X_i)$

¹Measured in terms of a *loss function*.

A random forest — properly tuned?

Predictive performance of an estimator can be measured in terms of some distance¹ between:

- 1) the observed outcome: Y_i
- 2) and the predicted conditional expectation: $\hat{f}_n(A_i, X_i)$

One example of a loss function $\mathcal{L}(f)(O)$ is the negative log-likelihood loss:

$$\mathcal{L}(\hat{f}_n)(Y_i, A_i, X_i) = -(Y_i \log(\hat{f}_n(A_i, X_i)) + (1 - Y_i) \log(1 - \hat{f}_n(A_i, X_i))).$$

¹Measured in terms of a *loss function*.

A random forest — properly tuned?

Predictive performance of an estimator can be measured in terms of some distance¹ between:

- 1) the observed outcome: Y_i
- 2) and the predicted conditional expectation: $\hat{f}_n(A_i, X_i)$

One example of a loss function $\mathcal{L}(f)(O)$ is the negative log-likelihood loss:

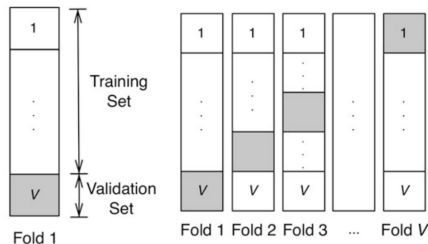
$$\mathcal{L}(\hat{f}_n)(Y_i, A_i, X_i) = -(Y_i \log(\hat{f}_n(A_i, X_i)) + (1 - Y_i) \log(1 - \hat{f}_n(A_i, X_i))).$$

The estimator \hat{f}_n closest to the true f_0 minimizes the risk:

$$\mathbb{E}_{P_0}[\mathcal{L}(\hat{f}_n)(Y_i, A_i, X_i)].$$

¹Measured in terms of a *loss function*.

A random forest — properly tuned?



The risk can be estimated in a cross-validation scheme.^a

I.e., for each sample split:

1. Each model is created and fitted on the training data: \hat{f}_n^{train} .
2. The quality of the model is checked on the validation data
 - ▶ Average of $\mathcal{L}(\hat{f}_n^{\text{train}})(O_i)$ in the validation sample.

^aTo measure performance on independent data.

A random forest — properly tuned?

Simulated example

- ▶ $X \sim \text{Unif}(-2, 2)$
- ▶ $X_1^{\text{noise}}, \dots, X_5^{\text{noise}} \sim N(0, 1)$
- ▶ $A \in \{0, 1\}$ with distribution given X given by:

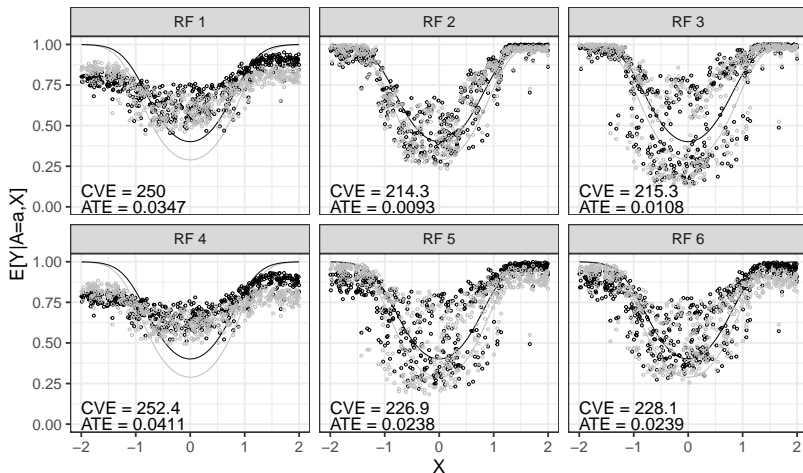
$$\text{logit } \mathbb{E}[A \mid X] = \gamma_0 + \gamma_X^\top X$$

- ▶ $Y \in \{0, 1\}$ with distribution given X and A given by:

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X^\top X^2$$


A random forest — properly tuned?

RF fitted with different values of tuning parameters (nodesize, mtry):



Different GLM models

GLM models fitted with different covariates and functional form of covariates:



A random forest — properly tuned?

This is all about constructing a good estimator for the conditional expectation f .

This does not necessarily translate into a good estimator for the target $\tilde{\Psi}(f, \mu_X)$.

A random forest — properly tuned?

This is all about constructing a good estimator for the conditional expectation f .

This does not necessarily translate into a good estimator for the target $\tilde{\Psi}(f, \mu_X)$.

TMLE is all about constructing a g-formula estimator which is a good estimator for *the target*.

Simulating simple data structure in R

Fix randomness:

```
set.seed(5)
```

Fix a sample size:

```
n <- 500
```

Generate covariate $X \in [-2, 2]$:

```
X <- runif(n, -2, 2)
```

Generate binary treatment decision A :

```
A <- rbinom(n, 1, prob=plogis(-0.25 + 1.2*X))
```

(corresponding to logit $\mathbb{E}[A | X] = \gamma_0 + \gamma_X X$)

Simulating simple data structure in R

Generate binary outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

Simulating simple data structure in R

Generate binary outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

First generate counterfactuals:

```
Y1 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*1))  
Y0 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*0))
```

Simulating simple data structure in R

Generate binary outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

First generate counterfactuals:

```
Y1 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*1))  
Y0 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*0))
```

We only observe the counterfactual outcome corresponding to the observed treatment level:

```
Y <- A*Y1 + (1-A)*Y0
```

Simulating simple data structure in R

Observed data:

```
      X A Y
1: -1.1991422 0 0
2:  0.7408744 1 0
3:  1.6675031 1 1
4: -0.8624022 0 1
5: -1.5813995 0 1
---
496: -0.3978523 1 0
497: -1.5069379 0 1
498:  1.8340120 1 1
499:  0.6349484 1 1
500: -0.5214807 0 1
```

Simulating simple data structure in R

Observed data:

	X	A	Y
1:	-1.1991422	0	0
2:	0.7408744	1	0
3:	1.6675031	1	1
4:	-0.8624022	0	1
5:	-1.5813995	0	1

496:	-0.3978523	1	0
497:	-1.5069379	0	1
498:	1.8340120	1	1
499:	0.6349484	1	1
500:	-0.5214807	0	1

Counterfactual data:

	X	Y1	Y0
1:	-1.1991422	0	1
2:	0.7408744	1	0
3:	1.6675031	1	1
4:	-0.8624022	0	1
5:	-1.5813995	1	1

496:	-0.3978523	0	1
497:	-1.5069379	0	1
498:	1.8340120	1	1
499:	0.6349484	0	0
500:	-0.5214807	0	0

Simulating simple data structure in R

Simulating many observations of counterfactuals allows us to approximate the **true ATE**:

```
X <- runif(1e6, -2, 2)
Y1 <- rbinom(1e6, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*1))
Y0 <- rbinom(1e6, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*0))
```

The **true ATE** is then approximately:

```
(true.ate <- mean(Y1 - Y0))
```

```
[1] 0.070292
```

since $ATE = \mathbb{E}_{P_0}[Y^1] - \mathbb{E}_{P_0}[Y^0]$.

Simulating simple data structure in R

Fit correctly specified parametric model:

```
fit.glm <- glm(Y~A+X.squared, data=dt[, X.squared:=X^2],  
              family=binomial)
```

Use model to estimate $f(1, X)$ for all subjects:

```
dt[, pred.glm.A1:=predict(fit.glm, type="response", newdata=  
  copy(dt)[, A:=1])]
```

And similarly $f(0, X)$ for all subjects:

```
dt[, pred.glm.A0:=predict(fit.glm, type="response", newdata=  
  copy(dt)[, A:=0])]
```

Then we can estimate the ATE by:

```
(fit.glm <- dt[, mean(pred.glm.A1-pred.glm.A0)])
```

```
[1] 0.04322891
```

Simulating simple data structure in R

Using a random forest (no tuning):

```
library(randomForestSRC)
fit.rf <- rfsrc(Y~A+X, data=dt)
dt[, pred.rf.A1:=predict(fit.rf, type="response", newdata=
  copy(dt)[, A:=1])$predicted]
dt[, pred.rf.A0:=predict(fit.rf, type="response", newdata=
  copy(dt)[, A:=0])$predicted]
(fit.rf <- dt[, mean(pred.rf.A1-pred.rf.A0)])
```

```
[1] 0.07005063
```

Simulating simple data structure in R

Using a misspecified parametric model:

```
fit.glm.mis <- glm(Y~A+X, data=dt, family=binomial)
dt[, pred.glm.mis.A1:=predict(fit.glm.mis, type="response",
                             newdata=copy(dt)[, A:=1])]
dt[, pred.glm.mis.A0:=predict(fit.glm.mis, type="response",
                             newdata=copy(dt)[, A:=0])]
(fit.glm.mis <- dt[, mean(pred.glm.mis.A1-pred.glm.mis.A0)])
```

[1] 0.09127889

Simulating simple data structure in R

We can investigate the properties of different estimators —

- ▶ We know the true value of ATE: $\psi_0 \approx 0.0702$
- ▶ We have generated the outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

- ▶ We have generated the treatment A according to

$$\text{logit } \mathbb{E}[A \mid X] = \gamma_0 + \gamma_X X$$

If we repeat the experiment of drawing n observations we would every time end up with a different realization of the particular estimator.

Different estimators

G-formula estimator Using an estimator \hat{f}_n for $f(a, X) = \mathbb{E}[Y \mid A = a, X]$, estimate the ATE by:

$$\hat{\psi}_n^{\text{g-formula}} = \frac{1}{n} \sum_{i=1}^n \{ \hat{f}_n(1, X_i) - \hat{f}_n(0, X_i) \}$$

Inverse probability weighted estimator Using an estimator $\hat{\pi}_n$ for $\pi(a \mid X) = P(A = a \mid X)$, estimate the ATE by:

$$\hat{\psi}_n^{\text{ipw}} = \frac{1}{n} \sum_{i=1}^n \left\{ \frac{A_i Y_i}{\hat{\pi}_n(A_i \mid X_i)} - \frac{(1 - A_i) Y_i}{\hat{\pi}_n(A_i \mid X_i)} \right\}$$

TMLE estimator Update the estimator $\hat{f}_n \mapsto \hat{f}_n^*$ in a "targeted way" using the information from the estimator $\hat{\pi}_n$, then estimate the ATE by:

$$\hat{\psi}_n^{\text{tmle}} = \frac{1}{n} \sum_{i=1}^n \{ \hat{f}_n^*(1, X_i) - \hat{f}_n^*(0, X_i) \}$$

Different estimators — `tmle` implementation

Today we will just (more or less blindly) use software to estimate TMLE.

Different estimators — tmle implementation

Today we will just (more or less blindly) use software to estimate TMLE.

```
library(tmle)
```

```
tmle(Y, A, X,  
      gform,  
      Qform,  
      SL.library,  
      family="binomial",  
      cvQinit=FALSE,  
      ...  
)
```

- ▶ $Y \in \mathbb{R}$ or $Y \in \{0, 1\}$
- ▶ $A \in \{0, 1\}$
- ▶ X a vector, matrix or a data frame

Different estimators — `tmle` implementation

- ▶ `gform`
 - ▶ optional regression formula for the propensity score π
 - ▶ on the form $A \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `Qform`
 - ▶ optional regression formula for the conditional expectation f
 - ▶ on the form $Y \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)

Different estimators — `tmle` implementation

- ▶ `gform`
 - ▶ optional regression formula for the propensity score π
 - ▶ on the form $A \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `Qform`
 - ▶ optional regression formula for the conditional expectation f
 - ▶ on the form $Y \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `cvQinit=FALSE`
 - ▶ default is `TRUE` which means cross-validated predicted values are estimated

Different estimators — `tmle` implementation

- ▶ `gform`
 - ▶ optional regression formula for the propensity score π
 - ▶ on the form $A \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `Qform`
 - ▶ optional regression formula for the conditional expectation f
 - ▶ on the form $Y \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `cvQinit=FALSE`
 - ▶ default is `TRUE` which means cross-validated predicted values are estimated
- ▶ `gbound`
 - ▶ truncation of predicted probabilities of treatment

Different estimators — `tmle` implementation

On a sidenote — tomorrow

- ▶ `Q.SL.library`
 - ▶ optional vector of prediction algorithms to use for SuperLearner in initial estimation of f
- ▶ `g.SL.library`
 - ▶ optional vector of prediction algorithms to use for SuperLearner in initial estimation of π
- ▶ `Q.discreteSL`
 - ▶ if TRUE, a discrete super learner is used (rather than ensemble)
 - ▶ default is FALSE
- ▶ `g.discreteSL`
 - ▶ if TRUE, a discrete super learner is used (rather than ensemble)
 - ▶ default is FALSE

Note: The discrete super learner simply picks an algorithm from its library by minimizing the cross-validated empirical risk with respect a loss function.

Different estimators — tmle implementation

What were the estimated IP weights?

```
summary(fit.tmle$g$g1W)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.04751	0.19441	0.49405	0.49400	0.79710	0.94109

Note that weights close to 0 or to 1 would indicate positivity issues.

Different estimators — tmle implementation

What were the estimated IP weights?

```
summary(fit.tmle$g$g1W)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.04751	0.19441	0.49405	0.49400	0.79710	0.94109

Note that weights close to 0 or to 1 would indicate positivity issues.

What truncation level was used?

```
fit.tmle$gbound
```

```
[1] 0.03598084 1.00000000
```

I.e., no weights were truncated.

Explorations based on simulated data

As part of the exercise we want to explore —

1. Comparing g-formula estimators for different estimators for f ; either different logistic regressions or different machine learning algorithms.
2. Properties of the g-formula estimator and the IP-weighted estimator, compared to the TMLE estimator.
3. Double robustness: Misspecification of the outcome regression (f).

The exercise is described in detail in: **day1-practical1.pdf**.

Nice asymptotic behaviour of an estimator

Asymptotic linearity

A very desirable property —

² $o_P(1)$ denotes a sequence which converges to zero in probability.

Asymptotic linearity

The empirical measure \mathbb{P}_n of the sample O_1, \dots, O_n :

$$\mathbb{P}_n h = \int h d\mathbb{P}_n = \frac{1}{n} \sum_{i=1}^n h(O_i).$$

A very desirable property —

An estimator $\hat{\psi}_n$ is \sqrt{n} -consistent and asymptotically linear with influence function $\phi(P_0)(O)$ if ²

$$\sqrt{n}(\hat{\psi}_n - \psi_0) = \sqrt{n}\mathbb{P}_n\phi(P_0) + o_P(1),$$

where $\mathbb{E}_{P_0}[\phi(P_0)(O)] = 0$ and $\mathbb{E}_{P_0}[\{\phi(P_0)(O)\}^2] < \infty$.

² $o_P(1)$ denotes a sequence which converges to zero in probability.

Asymptotic linearity

The empirical measure \mathbb{P}_n of the sample O_1, \dots, O_n :

$$\mathbb{P}_n h = \int h d\mathbb{P}_n = \frac{1}{n} \sum_{i=1}^n h(O_i).$$

A very desirable property —

An estimator $\hat{\psi}_n$ is \sqrt{n} -consistent and asymptotically linear with influence function $\phi(P_0)(O)$ if ²

$$\sqrt{n}(\hat{\psi}_n - \psi_0) = \sqrt{n}\mathbb{P}_n\phi(P_0) + o_P(1),$$

where $\mathbb{E}_{P_0}[\phi(P_0)(O)] = 0$ and $\mathbb{E}_{P_0}[\{\phi(P_0)(O)\}^2] < \infty$.

Then CLT + Slutsky implies:

$$\hat{\psi}_n \overset{as}{\sim} N(\Psi(P_0), \text{Var}(\phi(P_0))/n).$$

The estimator behaves asymptotically as an average of the influence function.

² $o_P(1)$ denotes a sequence which converges to zero in probability.

Asymptotic linearity

Simple example: Estimator for the mean $\psi_0 = \mathbb{E}[X]$:

$$\hat{\psi}_{n,0} = \frac{1}{n} \sum_{i=1}^n X_i$$

Then

$$\sqrt{n}(\hat{\psi}_n - \psi_0) = \sqrt{n} \frac{1}{n} \sum_{i=1}^n \underbrace{(X_i - \psi_0)}_{=\phi(P_0)(O_i)} = \sqrt{n} \mathbb{P}_n \phi(P_0)$$

$\hat{\psi}_{n,0}$ is linear and thus asymptotically linear.

Asymptotic linearity

Simple example: Estimator for the mean $\psi_0 = \mathbb{E}[X]$:

$$\hat{\psi}_{n,1} = \frac{1}{n} \sum_{i=1}^n X_i + \frac{1}{n}$$

Then

$$\sqrt{n}(\hat{\psi}_n - \psi_0) = \sqrt{n} \frac{1}{n} \sum_{i=1}^n \underbrace{(X_i - \psi_0)}_{=\phi(P_0)(O_i)} + \frac{\sqrt{n}}{n} = \sqrt{n} \mathbb{P}_n \phi(P_0) + \underbrace{\frac{1}{\sqrt{n}}}_{=o(1)}$$

$\hat{\psi}_{n,1}$ is asymptotically linear.

Asymptotic linearity

Simple example: Estimator for the mean $\psi_0 = \mathbb{E}[X]$:

$$\hat{\psi}_{n,2} = \frac{1}{n} \sum_{i=1}^n X_i + \frac{1}{n^{1/2+0.1}}$$

Then

$$\sqrt{n}(\hat{\psi}_n - \psi_0) = \sqrt{n} \frac{1}{n} \sum_{i=1}^n \underbrace{(X_i - \psi_0)}_{=\phi(P_0)(O_i)} + \frac{\sqrt{n}}{n^{1/2+0.1}} = \sqrt{n} \mathbb{P}_n \phi(P_0) + \underbrace{\frac{1}{n^{0.1}}}_{=o(1)}$$

$\hat{\psi}_{n,2}$ is asymptotically linear.

Asymptotic linearity

Simple example: Estimator for the mean $\psi_0 = \mathbb{E}[X]$:

$$\hat{\psi}_{n,3} = \frac{1}{n} \sum_{i=1}^n X_i + \frac{1}{n^{1/2-0.1}}$$

Then

$$\sqrt{n}(\hat{\psi}_n - \psi_0) = \sqrt{n} \frac{1}{n} \sum_{i=1}^n \underbrace{(X_i - \psi_0)}_{=\phi(P_0)(O_i)} + \frac{\sqrt{n}}{n^{1/2-0.1}} = \sqrt{n} \mathbb{P}_n \phi(P_0) + \underbrace{n^{0.1}}_{\rightarrow \infty}$$

$\hat{\psi}_{n,3}$ is **not** asymptotically linear.

Asymptotic linearity

An estimator $\hat{\psi}_n$ has **rate of convergence** $r_n \rightarrow \infty$ if ³

$$r_n(\hat{\psi}_n - \psi_0) = O_P(1), \quad \text{i.e.,} \quad \hat{\psi}_n - \psi_0 = O_P(1/r_n).$$

The convergence rate r_n tells us how fast $\hat{\psi}_n$ centers around ψ_0 , with the difference $\hat{\psi}_n - \psi_0$ behaving like $1/r_n$.

- ▶ One wants negligible bias such as to obtain reliable confidence intervals for ψ_0 .
- ▶ The bias of an asymptotically linear estimator converges to zero at a rate faster than $1/\sqrt{n}$.

Data-adaptive machine learning estimators rarely achieve this rate.

³ $O_P(1)$ denotes a sequence which is bounded in probability.

Asymptotic linearity

$$\sqrt{n}\hat{\psi}_{n,1} = \sqrt{n} \underbrace{\frac{1}{n} \sum_{i=1}^n X_i}_{\xrightarrow{P} \psi_0} + \underbrace{\frac{\sqrt{n}}{n}}_{\rightarrow 0}, \quad \text{i.e.,} \quad \sqrt{n}(\hat{\psi}_{n,1} - \psi_0) = o_P(1).$$

$$\sqrt{n}\hat{\psi}_{n,2} = \sqrt{n} \underbrace{\frac{1}{n} \sum_{i=1}^n X_i}_{\xrightarrow{P} \psi_0} + \underbrace{\frac{\sqrt{n}}{n^{1/2+0.1}}}_{\rightarrow 0}, \quad \text{i.e.,} \quad \sqrt{n}(\hat{\psi}_{n,2} - \psi_0) = o_P(1).$$

$$\sqrt{n}\hat{\psi}_{n,3} = \sqrt{n} \underbrace{\frac{1}{n} \sum_{i=1}^n X_i}_{\xrightarrow{P} \psi_0} + \underbrace{\frac{\sqrt{n}}{n^{1/2-0.1}}}_{\rightarrow \infty}, \quad \text{i.e.,} \quad \sqrt{n}(\hat{\psi}_{n,3} - \psi_0) \xrightarrow{P} \infty.$$

Asymptotic linearity

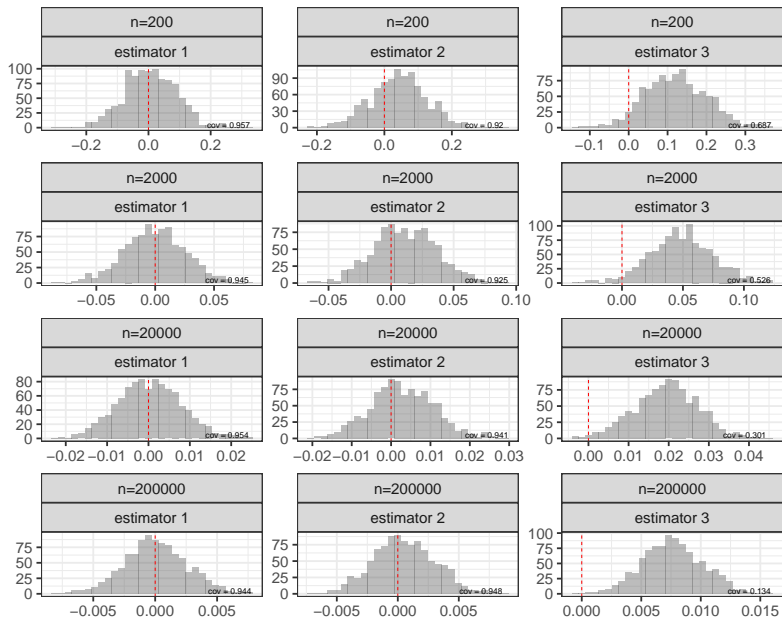
$$\sqrt{n}\hat{\psi}_{n,1} = \underbrace{\sqrt{n} \frac{1}{n} \sum_{i=1}^n X_i}_{\overset{P}{\rightarrow} \psi_0} + \underbrace{\frac{\sqrt{n}}{n}}_{\rightarrow 0}, \quad \text{i.e.,} \quad \sqrt{n}(\hat{\psi}_{n,1} - \psi_0) = o_P(1).$$

$$\sqrt{n}\hat{\psi}_{n,2} = \underbrace{\sqrt{n} \frac{1}{n} \sum_{i=1}^n X_i}_{\overset{P}{\rightarrow} \psi_0} + \underbrace{\frac{\sqrt{n}}{n^{1/2+0.1}}}_{\rightarrow 0}, \quad \text{i.e.,} \quad \sqrt{n}(\hat{\psi}_{n,2} - \psi_0) = o_P(1).$$

$$\sqrt{n}\hat{\psi}_{n,3} = \underbrace{\sqrt{n} \frac{1}{n} \sum_{i=1}^n X_i}_{\overset{P}{\rightarrow} \psi_0} + \underbrace{\frac{\sqrt{n}}{n^{1/2-0.1}}}_{\rightarrow \infty}, \quad \text{i.e.,} \quad \sqrt{n}(\hat{\psi}_{n,3} - \psi_0) \overset{P}{\rightarrow} \infty.$$

[The remainder term that determines the asymptotic bias the estimator].

Asymptotic linearity



Asymptotic linearity

A key component in constructing a \sqrt{n} -consistent and asymptotically linear estimator, *even when using machine learning estimation*, is the so-called **the efficient influence function** (also known as the canonical gradient).

Asymptotic linearity

A key component in constructing a \sqrt{n} -consistent and asymptotically linear estimator, *even when using machine learning estimation*, is the so-called **the efficient influence function** (also known as the canonical gradient).

- ▶ The efficient influence function provides a nonparametric lower bound for the estimation problem.
- ▶ Tells us how to do bias-correction.
- ▶ With the bias-correction, the remainder term that we need to control to have \sqrt{n} -consistency and asymptotic linearity admits a nice structure that we *can* control.

⇒ TMLE.