

Day 1, Lecture 3


Estimating the target

Estimating the target

In this lecture, our goal is to:

1. Describe simple causal inference estimation methods for the average treatment effect, including the estimation equations estimator, an inverse probability weighted estimator and the g-formula based estimator.
2. Compare the assumptions and performance of “double robust” estimation methods to estimation methods based on inverse probability weighting and (g-formula based) standardization.

Estimating the target



- ▶ One estimator is not more causal than another.
- ▶ Different estimators are based on different nuisance parameters and have different statistical properties (bias/variance).

G-formula versus IP-weighting

- G-formula
1. Estimate nuisance parameters
 $f(a, x) = \mathbb{E}[Y \mid A = a, X = x]$ and the average over the marginal distribution μ_X of X
 2. Plug in to estimate the ATE:

$$\hat{\psi}_n^{\text{g-formula}} = \tilde{\Psi}(\hat{f}_n, \hat{\mu}_X) = \int_{\mathbb{R}^d} (\hat{f}_n(1, x) - \hat{f}_n(0, x)) d\hat{\mu}_X(x)$$

- IP-weighting
1. Estimate nuisance parameters
 $\pi(a \mid x) = P(A = a \mid X = x)$ and the average over the distribution P of O
 2. Plug in to estimate the ATE:

$$\hat{\psi}_n^{\text{ipw}} = \tilde{\Psi}_{\text{ipw}}(\hat{\pi}_n, \hat{P}_n) = \int_{\mathbb{R}^d} \sum_{a=0,1} \sum_{y=0,1} \left(\frac{ay}{\hat{\pi}_n(a \mid x)} - \frac{(1-a)y}{\hat{\pi}_n(a \mid x)} \right) d\hat{P}_n(x)$$

Estimating equation (EE) estimator

EE estimator 1. Estimate nuisance parameters

$f(a, x) = \mathbb{E}[Y \mid A = a, X = x]$, $\pi(a \mid x) = \mathbb{E}[A \mid X = x]$
and the average over the distribution P of O

2. Plug in to estimate the ATE:

$$\hat{\psi}_n^{\text{ee}} = \tilde{\Psi}_{\text{ee}}(\hat{f}_n, \hat{\pi}_n, \hat{P}_n) = \int_{\mathbb{R}^d} \sum_{a=0,1} \sum_{y=0,1} \left\{ \left(\frac{a}{\hat{\pi}_n(a \mid x)} - \frac{1-a}{\hat{\pi}_n(a \mid x)} \right) (y - \hat{f}_n(a, x)) \right. \\ \left. + \hat{f}_n(1, x) - \hat{f}_n(0, x) \right\} d\hat{P}_n(o)$$

G-formula versus IP-weighting versus EE/TMLE

Estimation of the averages over μ_X and P is straightforward using the empirical average over the observed data.

This yields:

G-formula estimator:
$$\hat{\psi}_n^{\text{g-formula}} = \frac{1}{n} \sum_{i=1}^n \{ \hat{f}_n(1, X_i) - \hat{f}_n(0, X_i) \}$$

IP-weighted estimator:
$$\hat{\psi}_n^{\text{ipw}} = \frac{1}{n} \sum_{i=1}^n \left\{ \frac{A_i Y_i}{\hat{\pi}_n(A_i | X_i)} - \frac{(1 - A_i) Y_i}{\hat{\pi}_n(A_i | X_i)} \right\}$$

EE estimator:
$$\hat{\psi}_n^{\text{ee}} = \frac{1}{n} \sum_{i=1}^n \left\{ \left(\frac{A_i}{\hat{\pi}_n(A_i | X_i)} - \frac{(1 - A_i)}{\hat{\pi}_n(A_i | X_i)} \right) (Y_i - \hat{f}_n(A_i, X_i)) + \hat{f}_n(1, X_i) - \hat{f}_n(0, X_i) \right\}.$$

G-formula versus IP-weighting versus EE/TMLE

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

¹which we get back to.

G-formula versus IP-weighting versus EE/TMLE

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

EE estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

¹which we get back to.

G-formula versus IP-weighting versus EE/TMLE

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

EE estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

- ▶ consistent if either \hat{f}_n or $\hat{\pi}_n$ is consistent (commonly known as "double robustness").

¹which we get back to.

G-formula versus IP-weighting versus EE/TMLE

G-formula estimator requires estimator \hat{f}_n for conditional expectation f .

- ▶ consistent if \hat{f}_n is consistent.

IP-weighted estimator requires estimator $\hat{\pi}_n$ for the propensity score π .

- ▶ consistent if $\hat{\pi}_n$ is consistent.

EE estimator requires estimators \hat{f}_n and $\hat{\pi}_n$ for conditional expectation f and propensity score π .

- ▶ consistent if either \hat{f}_n or $\hat{\pi}_n$ is consistent (commonly known as "double robustness").
- ▶ the EE estimator and the TMLE estimator share the same large-sample properties,¹ and particularly this property.

¹which we get back to.

G-formula versus IP-weighting versus EE/TMLE

"Double robustness" —

SMALL EXERCISE:

By the law of large numbers, the EE estimator converges in probability to:

$$\mathbb{E}_{P_0} \left[\left(\frac{A}{\pi(A|X)} - \frac{1-A}{\pi(A|X)} \right) (Y - f(A, X)) + f(1, X) - f(0, X) \right] \quad (1)$$

where (f, π) denotes the limit of $(\hat{f}_n, \hat{\pi}_n)$. Compute (1) when:

1. $f = f_0$ (i.e., the outcome regression is consistently estimated), and
2. $\pi = \pi_0$ (i.e., the propensity score is consistently estimated).

G-formula versus IP-weighting versus EE/TMLE

Can't we just construct a good g-formula estimator???

G-formula versus IP-weighting versus EE/TMLE

Can't we just construct a good g-formula estimator???

- ▶ a logistic regression — great if correctly specified, but horrible if not.
- ▶ a random forest — properly tuned?

A random forest — properly tuned?

Predictive performance of an estimator can be measured in terms of some distance² between:

- 1) the observed outcome: Y_i
- 2) and the predicted conditional expectation: $\hat{f}_n(A_i, X_i)$

²Measured in terms of a *loss function*.

A random forest — properly tuned?

Predictive performance of an estimator can be measured in terms of some distance² between:

- 1) the observed outcome: Y_i
- 2) and the predicted conditional expectation: $\hat{f}_n(A_i, X_i)$

One example of a loss function $\mathcal{L}(f)(O)$ is the negative log-likelihood loss:

$$\mathcal{L}(\hat{f}_n)(Y_i, A_i, X_i) = -(Y_i \log(\hat{f}_n(A_i, X_i)) + (1 - Y_i) \log(1 - \hat{f}_n(A_i, X_i))).$$

²Measured in terms of a *loss function*.

A random forest — properly tuned?

Predictive performance of an estimator can be measured in terms of some distance² between:

- 1) the observed outcome: Y_i
- 2) and the predicted conditional expectation: $\hat{f}_n(A_i, X_i)$

One example of a loss function $\mathcal{L}(f)(O)$ is the negative log-likelihood loss:

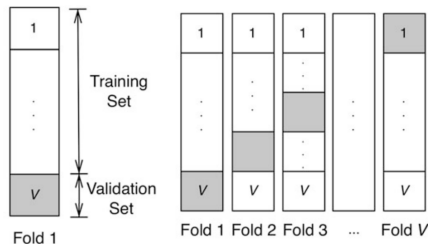
$$\mathcal{L}(\hat{f}_n)(Y_i, A_i, X_i) = -(Y_i \log(\hat{f}_n(A_i, X_i)) + (1 - Y_i) \log(1 - \hat{f}_n(A_i, X_i))).$$

The estimator \hat{f}_n closest to the true f_0 minimizes the risk:

$$\mathbb{E}_{P_0}[\mathcal{L}(\hat{f}_n)(Y_i, A_i, X_i)].$$

²Measured in terms of a *loss function*.

A random forest — properly tuned?



The risk can be estimated in a cross-validation scheme.^a

I.e., for each sample split:

1. Each model is created and fitted on the training data: \hat{f}_n^{train} .
2. The quality of the model is checked on the validation data
 - ▶ Average of $\mathcal{L}(\hat{f}_n^{\text{train}})(O_i)$ in the validation sample.

^aTo measure performance on independent data.

A random forest — properly tuned?

The **loss-based cross-validated error** can then be obtained for each \hat{f}^j in a "library" of learners $\{\hat{f}_n^1, \hat{f}_n^2, \dots, \hat{f}_n^J\}$ as

$$\frac{1}{V} \sum_{v=1}^V \frac{1}{\#_v} \sum_{i \in v} \mathcal{L}(\hat{f}_{-v}^j)(O_i) \quad (2)$$

We can for example use this to pick the estimator \hat{f}^{j*} which minimizes (2).

- ▶ With a library of random forests, with different values of hyperparameters, for example, this is just a version of tuning.
- ▶ With a library of parametric regression models, this is a form of model selection.

Weighting (or picking out) the models/algorithms according to how well they fit the data, with performance measured in a cross-validation scheme, is what we call **super (ensemble) learning**.

A random forest — properly tuned?

Simulated example

- ▶ $X \sim \text{Unif}(-2, 2)$
- ▶ $X_1^{\text{noise}}, \dots, X_5^{\text{noise}} \sim N(0, 1)$
- ▶ $A \in \{0, 1\}$ with distribution given X given by:


$$\text{logit } \mathbb{E}[A \mid X] = \gamma_0 + \gamma_X^\top X$$

- ▶ $Y \in \{0, 1\}$ with distribution given X and A given by:

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X^\top X^2$$

A random forest — properly tuned?

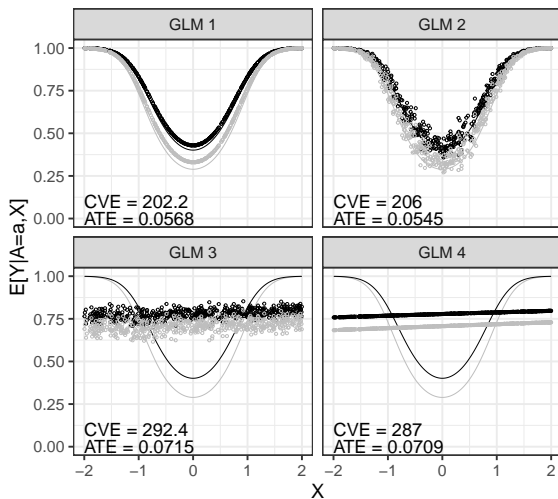
RF fitted with different values of tuning parameters (nodesize, mtry):



value of a → 0 → 1

Different GLM models

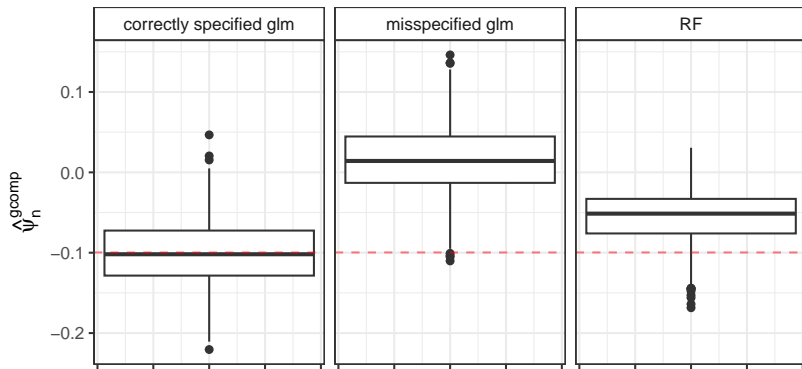
GLM models fitted with different covariates and functional form of covariates:



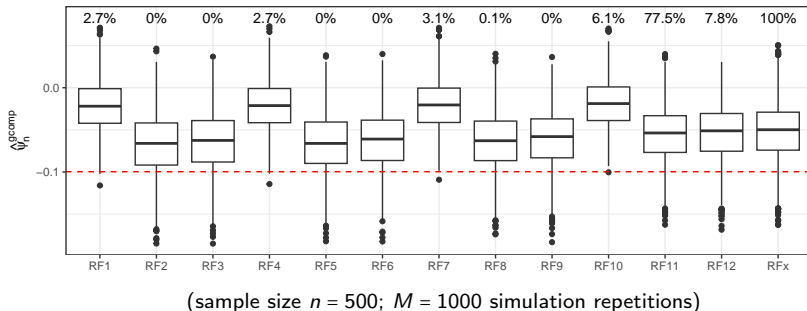
A random forest — properly tuned?

ATE g-computation estimation in simulated setting:

(sample size $n = 500$; $M = 1000$ simulation repetitions)



A random forest — properly tuned?



Furthest to the right (marked **RFx**) is the sampling distribution of the **direct ATE g-computation estimates** obtained using the forest with the lowest cross-validated error in each simulation repetition.

A random forest — properly tuned?

Finding the estimator \hat{f}_n minimizing the cross-validated empirical risk is all about constructing a good estimator for the conditional expectation f .

This does not necessarily translate into a good estimator for the target estimand $\Psi(P)$.

A random forest — properly tuned?

Finding the estimator \hat{f}_n minimizing the cross-validated empirical risk is all about constructing a good estimator for the conditional expectation f .

This does not necessarily translate into a good estimator for the target estimand $\Psi(P)$.

TMLE is all about constructing a g-formula estimator which is a good estimator for *the target*.

Simulating simple data structure in R

Fix randomness:

```
set.seed(5)
```

Fix a sample size:

```
n <- 500
```

Generate covariate $X \in [-2, 2]$:

```
X <- runif(n, -2, 2)
```

Generate binary treatment decision A :

```
A <- rbinom(n, 1, prob=plogis(-0.25 + 1.2*X))
```

(corresponding to logit $\mathbb{E}[A | X] = \gamma_0 + \gamma_X X$)

Simulating simple data structure in R

Generate binary outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

Simulating simple data structure in R

Generate binary outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

First generate counterfactuals:

```
Y1 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*1))  
Y0 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*0))
```

Simulating simple data structure in R

Generate binary outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

First generate counterfactuals:

```
Y1 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*1))  
Y0 <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*0))
```

We only observe the counterfactual outcome corresponding to the observed treatment level:

```
Y <- A*Y1 + (1-A)*Y0
```

Simulating simple data structure in R

Observed data:

```
      X A Y
1: -1.1991422 0 0
2:  0.7408744 1 0
3:  1.6675031 1 1
4: -0.8624022 0 1
5: -1.5813995 0 1
---
496: -0.3978523 1 0
497: -1.5069379 0 1
498:  1.8340120 1 1
499:  0.6349484 1 1
500: -0.5214807 0 1
```

Simulating simple data structure in R

Observed data:

	X	A	Y
1:	-1.1991422	0	0
2:	0.7408744	1	0
3:	1.6675031	1	1
4:	-0.8624022	0	1
5:	-1.5813995	0	1

496:	-0.3978523	1	0
497:	-1.5069379	0	1
498:	1.8340120	1	1
499:	0.6349484	1	1
500:	-0.5214807	0	1

Counterfactual data:

	X	Y1	Y0
1:	-1.1991422	0	1
2:	0.7408744	1	0
3:	1.6675031	1	1
4:	-0.8624022	0	1
5:	-1.5813995	1	1

496:	-0.3978523	0	1
497:	-1.5069379	0	1
498:	1.8340120	1	1
499:	0.6349484	0	0
500:	-0.5214807	0	0

Simulating simple data structure in R

Simulating many observations of counterfactuals allows us to approximate the **true ATE**:

```
X <- runif(1e6, -2, 2)
Y1 <- rbinom(1e6, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*1))
Y0 <- rbinom(1e6, 1, prob=plogis(-0.9 + 1.9*X^2 + 0.5*0))
```

The **true ATE** is then approximately:

```
(true.ate <- mean(Y1 - Y0))
```

```
[1] 0.070292
```

since $ATE = \mathbb{E}_{P_0}[Y^1] - \mathbb{E}_{P_0}[Y^0]$.

Simulating simple data structure in R

Fit correctly specified parametric model:

```
fit.glm <- glm(Y~A+X.squared, data=dt[, X.squared:=X^2],  
              family=binomial)
```

Use model to estimate $f(1, X)$ for all subjects:

```
dt[, pred.glm.A1:=predict(fit.glm, type="response", newdata=  
  copy(dt)[, A:=1])]
```

And similarly $f(0, X)$ for all subjects:

```
dt[, pred.glm.A0:=predict(fit.glm, type="response", newdata=  
  copy(dt)[, A:=0])]
```

Then we can estimate the ATE by:

```
(fit.glm <- dt[, mean(pred.glm.A1-pred.glm.A0)])
```

```
[1] 0.04322891
```

Simulating simple data structure in R

Using a random forest (no tuning):

```
library(randomForestSRC)
fit.rf <- rfsrc(Y~A+X, data=dt)
dt[, pred.rf.A1:=predict(fit.rf, type="response", newdata=
  copy(dt)[, A:=1])$predicted]
dt[, pred.rf.A0:=predict(fit.rf, type="response", newdata=
  copy(dt)[, A:=0])$predicted]
(fit.rf <- dt[, mean(pred.rf.A1-pred.rf.A0)])
```

```
[1] 0.07005063
```

Simulating simple data structure in R

Using a misspecified parametric model:

```
fit.glm.mis <- glm(Y~A+X, data=dt, family=binomial)
dt[, pred.glm.mis.A1:=predict(fit.glm.mis, type="response",
                             newdata=copy(dt)[, A:=1])]
dt[, pred.glm.mis.A0:=predict(fit.glm.mis, type="response",
                             newdata=copy(dt)[, A:=0])]
(fit.glm.mis <- dt[, mean(pred.glm.mis.A1-pred.glm.mis.A0)])
```

[1] 0.09127889

Simulating simple data structure in R

We can investigate the properties of different estimators —

- ▶ We know the true value of ATE: $\psi_0 \approx 0.0702$
- ▶ We have generated the outcome Y according to

$$\text{logit } \mathbb{E}[Y \mid A, X] = \beta_0 + \beta_A A + \beta_X X^2$$

- ▶ We have generated the treatment A according to

$$\text{logit } \mathbb{E}[A \mid X] = \gamma_0 + \gamma_X X$$

If we repeat the experiment of drawing n observations we would every time end up with a different realization of the particular estimator.

Different estimators

G-formula estimator Using an estimator \hat{f}_n for $f(a, X) = \mathbb{E}[Y \mid A = a, X]$, estimate the ATE by:

$$\hat{\psi}_n^{\text{g-formula}} = \frac{1}{n} \sum_{i=1}^n \{ \hat{f}_n(1, X_i) - \hat{f}_n(0, X_i) \}$$

Inverse probability weighted estimator Using an estimator $\hat{\pi}_n$ for $\pi(a \mid X) = P(A = a \mid X)$, estimate the ATE by:

$$\hat{\psi}_n^{\text{ipw}} = \frac{1}{n} \sum_{i=1}^n \left\{ \frac{A_i Y_i}{\hat{\pi}_n(A_i \mid X_i)} - \frac{(1 - A_i) Y_i}{\hat{\pi}_n(A_i \mid X_i)} \right\}$$

Different estimators

EE estimator Using an estimator \hat{f}_n for $f(a, X) = \mathbb{E}[Y \mid A = a, X]$ and an estimator $\hat{\pi}_n$ for $\pi(a \mid X) = P(A = a \mid X)$, estimate the ATE by:

$$\hat{\psi}_n^{\text{ee}} = \frac{1}{n} \sum_{i=1}^n \left\{ \frac{A_i}{\hat{\pi}_n(A_i \mid X_i)} - \frac{(1 - A_i)}{\hat{\pi}_n(A_i \mid X_i)} (Y_i - \hat{f}_n(A_i, X_i)) \right. \\ \left. + \hat{f}_n(1, X_i) - \hat{f}_n(0, X_i) \right\}$$

TMLE estimator Update the estimator $\hat{f}_n \mapsto \hat{f}_n^*$ in a "targeted way" using the information from the estimator $\hat{\pi}_n$, then estimate the ATE by:

$$\hat{\psi}_n^{\text{tmle}} = \frac{1}{n} \sum_{i=1}^n \{ \hat{f}_n^*(1, X_i) - \hat{f}_n^*(0, X_i) \}$$

Different estimators — `tmle` implementation

Today we will just (more or less blindly) use software to use TMLE.

Different estimators — tmle implementation

Today we will just (more or less blindly) use software to use TMLE.

```
library(tmle)
```

```
tmle(Y, A, X,  
      gform,  
      Qform,  
      SL.library,  
      family="binomial",  
      cvQinit=FALSE,  
      ...  
)
```

- ▶ $Y \in \mathbb{R}$ or $Y \in \{0, 1\}$
- ▶ $A \in \{0, 1\}$
- ▶ X a vector, matrix or a data frame

Different estimators — `tmle` implementation

- ▶ `gform`
 - ▶ optional regression formula for the propensity score π
 - ▶ on the form $A \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `Qform`
 - ▶ optional regression formula for the conditional expectation f
 - ▶ on the form $Y \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)

Different estimators — `tmle` implementation

- ▶ `gform`
 - ▶ optional regression formula for the propensity score π
 - ▶ on the form $A \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `Qform`
 - ▶ optional regression formula for the conditional expectation f
 - ▶ on the form $Y \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `cvQinit=FALSE`
 - ▶ default is `TRUE` which means cross-validated predicted values are estimated

Different estimators — `tmle` implementation

- ▶ `gform`
 - ▶ optional regression formula for the propensity score π
 - ▶ on the form $A \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `Qform`
 - ▶ optional regression formula for the conditional expectation f
 - ▶ on the form $Y \sim X_1 + X_2$
 - ▶ (overrides call to `SuperLearner`)
- ▶ `cvQinit=FALSE`
 - ▶ default is `TRUE` which means cross-validated predicted values are estimated
- ▶ `gbound`
 - ▶ truncation of predicted probabilities of treatment

Different estimators — `tmle` implementation

On a sidenote (we return to this)

- ▶ `Q.SL.library`
 - ▶ optional vector of prediction algorithms to use for SuperLearner in initial estimation of f
- ▶ `g.SL.library`
 - ▶ optional vector of prediction algorithms to use for SuperLearner in initial estimation of π
- ▶ `Q.discreteSL`
 - ▶ if TRUE, a discrete super learner is used (rather than ensemble)
 - ▶ default is FALSE
- ▶ `g.discreteSL`
 - ▶ if TRUE, a discrete super learner is used (rather than ensemble)
 - ▶ default is FALSE

Note: The discrete super learner simply picks an algorithm from its library by minimizing the cross-validated empirical risk with respect a loss function.

Different estimators — tmle implementation

What were the estimated IP weights?

```
summary(fit.tmle$g$g1W)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.04751	0.19441	0.49405	0.49400	0.79710	0.94109

Note that weights close to 0 or to 1 would indicate positivity issues.

Different estimators — tmle implementation

What were the estimated IP weights?

```
summary(fit.tmle$g$g1W)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.04751	0.19441	0.49405	0.49400	0.79710	0.94109

Note that weights close to 0 or to 1 would indicate positivity issues.

What truncation level was used?

```
fit.tmle$gbound
```

```
[1] 0.03598084 1.00000000
```

I.e., no weights were truncated.

Practical 1: Explorations based on simulated data

As part of the exercise we will explore —

1. Comparing g-formula estimators for different estimators for f ; either different logistic regressions or different machine learning algorithms.
2. Properties of the g-formula estimator and the IP-weighted estimator, compared to the TMLE estimator.
3. Double robustness: Misspecification of the outcome regression (f).

The exercise is described in detail in: **day1-practical1.pdf**.