

Day 1, Practical 1, Hely's solution

Helene Charlotte Wiese Rytgaard

September 28, 2021

1 Simulating data

Task 1.

```
library(data.table)

sim.fun <- function(n) {

  X1 <- runif(n, -2, 2)
  X2 <- rnorm(n, 0, 0.2)
  X3 <- rbinom(n, 1, prob=0.2)
  A <- rbinom(n, 1, prob=plogis(-0.25+0.8*X1+0.25*X3))
  Y <- rbinom(n, 1, prob=plogis(-0.9+1.9*X1^2+0.6*X2+0.5*A))

  return(data.table(X1=X1, X2=X2, X3=X3, A=A, Y=Y))

}
```

Task 2.

```
library(data.table)

sim.fun <- function(n, a=NULL) {
  X1 <- runif(n, -2, 2)
  X2 <- rnorm(n)
  X3 <- rbinom(n, 1, 0.2)
  if (length(a)>0) {
    A <- a
  } else {
    A <- rbinom(n, 1, prob=plogis(-0.25 + 0.8*X1 + 0.25*X3))
  }
  Y <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X1^2 + 0.6*X2 + 0.5*A))
  if (length(a)>0) {
    return(mean(Y))
  } else {
    return(data.table(id=1:n,X1=X1,X2=X2,X3=X3,A=A,Y=Y))
  }
}
```

```
set.seed(12)
message(paste0("EY1 = ", E.Y1 <- sim.fun(1e6, a=1)))
message(paste0("EY0 = ", E.Y0 <- sim.fun(1e6, a=0)))
message(paste0("ATE = ", ATE <- E.Y1 - E.Y0))
```

```
EY1 = 0.749921
EY0 = 0.68208
ATE = 0.0678409999999999
```

2 Estimation

Task 3.

```
set.seed(15)
head(sim.data <- sim.fun(1000))
```

```
id      X1      X2 X3 A Y
1:  1  0.4084562  0.38996075  0 0 0
2:  2 -1.2198243 -1.67449303  1 0 0
3:  3  1.8658349 -2.22881407  0 1 1
4:  4  0.6036221 -0.01388672  0 0 0
5:  5 -0.5317124  0.57686435  0 0 0
6:  6  1.9554368  0.15718650  0 0 1
```

```
message("fitted model for the outcome regression:")
summary(fit.f <- glm(Y~A+X1+X2+X3, family=binomial, data=sim.data))
message("-----")
message("fitted model for the propensity score:")
summary(fit.pi <- glm(A~X1+X2+X3, family=binomial, data=sim.data))
```

fitted model for the outcome regression:

Call:

```
glm(formula = Y ~ A + X1 + X2 + X3, family = binomial, data = sim.data)
```

Deviance Residuals:

```
      Min       1Q   Median       3Q      Max
-2.2530  -1.2483   0.6745   0.8255   1.3174
```

Coefficients:

```
              Estimate Std. Error z value      Pr(>|z|)
(Intercept)  0.76917     0.10360   7.424 0.000000000000113 ***
A             0.46609     0.16396   2.843   0.00447 **
X1            -0.01980     0.06944  -0.285   0.77554
X2             0.45090     0.07413   6.083 0.000000001180179 ***
X3            0.23320     0.19117   1.220   0.22253
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 1170.5  on 999  degrees of freedom
Residual deviance: 1121.4  on 995  degrees of freedom
AIC: 1131.4
```

Number of Fisher Scoring iterations: 4

fitted model for the propensity score:

Call:

```
glm(formula = A ~ X1 + X2 + X3, family = binomial, data = sim.data)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.9285	-0.9355	-0.5552	0.9638	2.0656

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.314837	0.079178	-3.976	0.00007 ***
X1	0.843342	0.066527	12.677	< 2e-16 ***
X2	-0.004127	0.068828	-0.060	0.9522
X3	0.379385	0.179716	2.111	0.0348 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1380.8 on 999 degrees of freedom
Residual deviance: 1182.2 on 996 degrees of freedom
AIC: 1190.2

Number of Fisher Scoring iterations: 4

```
##-- g-formula;  
sim.data[, pred.EY1:=predict(fit.f, type="response", newdata=copy(sim.data)[, A:=1])]  
sim.data[, pred.EY0:=predict(fit.f, type="response", newdata=copy(sim.data)[, A:=0])]  
message(paste0("g-formula estimate = ", round(sim.data[, mean(pred.EY1 - pred.EY0)],  
5)))  
##-- ipw;  
sim.data[, pred.pi1:=predict(fit.pi, type="response", newdata=sim.data)]  
message(paste0("ipw estimate = ", round(sim.data[, mean(A*Y/pred.pi1 - (1-A)*Y/(1-pred  
.pi1))], 5)))
```

g-formula estimate = 0.08732

ipw estimate = 0.05979

Task 4.

```
fit.f2 <- glm(Y~A+X1.squared+X2+X3, family=binomial, data=sim.data[, X1.squared:=X1  
^2])  
fit.pi <- glm(A~X1+X2+X3, family=binomial, data=sim.data)  
##-- g-formula;  
sim.data[, pred.EY1:=predict(fit.f2, type="response", newdata=copy(sim.data)[, A:=1])]  
sim.data[, pred.EY0:=predict(fit.f2, type="response", newdata=copy(sim.data)[, A:=0])]  
message(paste0("g-formula estimate = ", round(sim.data[, mean(pred.EY1 - pred.EY0)],  
5)))  
##-- ipw;
```

```
sim.data[, pred.pi1:=predict(fit.pi, type="response", newdata=sim.data)]
message(paste0("ipw estimate = ", round(sim.data[, mean(A*Y/pred.pi1 - (1-A)*Y/(1-pred
.pi1))], 5)))
```

g-formula estimate = 0.06703
ipw estimate = 0.05979

Task 5.

```
library(randomForestSRC)
fit.rf <- rfsrc(Y~A+X1+X2+X3, data=sim.data)
##-- g-formula;
sim.data[, pred.EY1:=predict(fit.rf, type="response", newdata=copy(sim.data)[, A:=1])$
predicted]
sim.data[, pred.EY0:=predict(fit.rf, type="response", newdata=copy(sim.data)[, A:=0])$
predicted]
message(paste0("g-formula estimate (RF) = ", round(sim.data[, mean(pred.EY1 - pred.EY0
)], 5)))
```

g-formula estimate (RF) = 0.06118

Task 6.

```
library(tml)
tmle.fit <- tmle(Y=sim.data$Y, A=sim.data$A,
  cbind(X1=sim.data$X1,X2=sim.data$X2,X3=sim.data$X3),
  gform=A~X1+X2+X3, ## treatment model
  Qform=Y~A+X1+X2+X3, ## outcome model
  family="binomial",
  cvQinit=FALSE)
##-- get the ATE estimate:
tmle.fit$estimates$ATE$psi
```

[1] 0.06626344

```
tmle.fit$Qinit$coef
fit.f$coef
```

(Intercept)	A	X1	X2	X3
0.76917331	0.46609109	-0.01979933	0.45089688	0.23320073
(Intercept)	A	X1	X2	X3
0.76917331	0.46609109	-0.01979933	0.45089688	0.23320073

Task 7.

```
tmle.fit2 <- tmle(Y=sim.data$Y, A=sim.data$A,
  cbind(X1=sim.data$X1, X1.squared=sim.data$X1^2,X2=sim.data$X2,X3=sim.data$X3
),
  gform=A~X1+X2+X3, ## treatment model
  Qform=Y~A+X1.squared+X2+X3, ## outcome model
  family="binomial",
  cvQinit=FALSE)
##-- get the ATE estimate:
tmle.fit2$estimates$ATE$psi
```

[1] 0.06791028

```
tmle.fit2$Qinit$coef  
fit.f2$coef
```

(Intercept)	A	X1.squared	X2	X3
-0.8362875	0.5140438	1.8457299	0.5762989	0.1435289

(Intercept)	A	X1.squared	X2	X3
-0.8362875	0.5140438	1.8457299	0.5762989	0.1435289

3 Simulation study

Task 8.

```
fit.g.glm1 <- list()  
fit.g.glm2 <- list()  
fit.g.rf <- list()  
fit.ipw <- list()  
fit.tmle <- list()  
fit.one.step <- list()  
  
for (m in 1:500) {  
  
  set.seed(m+110)  
  sim.data <- sim.fun(1000)  
  
  fit.f <- glm(Y~A+X1+X2+X3, family=binomial, data=sim.data)  
  fit.f2 <- glm(Y~A+X1.squared+X2+X3, family=binomial, data=sim.data[, X1.squared:=  
    X1^2])  
  fit.rf <- rfsrc(Y~A+X1+X2+X3, data=sim.data)  
  fit.pi <- glm(A~X1+X2+X3, family=binomial, data=sim.data)  
  
  ##-- g-formula (section 3.1);  
  sim.data[, pred.f1:=predict(fit.f, type="response", newdata=copy(sim.data)[, A  
    :=1])]   
  sim.data[, pred.f0:=predict(fit.f, type="response", newdata=copy(sim.data)[, A  
    :=0])]   
  fit.g.glm1[[m]] <- sim.data[, mean(pred.f1 - pred.f0)]  
  
  ##-- g-formula (section 3.2);  
  sim.data[, pred.f1:=predict(fit.f2, type="response", newdata=copy(sim.data)[, A  
    :=1])]   
  sim.data[, pred.f0:=predict(fit.f2, type="response", newdata=copy(sim.data)[, A  
    :=0])]   
  fit.g.glm2[[m]] <- sim.data[, mean(pred.f1 - pred.f0)]  
  
  ##-- g-formula based on RF (section 3.3);  
  sim.data[, pred.f1:=predict(fit.rf, type="response", newdata=copy(sim.data)[, A  
    :=1])$predicted]   
  sim.data[, pred.f0:=predict(fit.rf, type="response", newdata=copy(sim.data)[, A  
    :=0])$predicted]   
  fit.g.rf[[m]] <- sim.data[, mean(pred.f1 - pred.f0)]  
}
```

```

##-- ipw (section 3.1);
sim.data[, pred.pi1:=predict(fit.pi, type="response", newdata=sim.data)]
fit.ipw[[m]] <- sim.data[, mean(A*Y/pred.pi1 - (1-A)*Y/(1-pred.pi1))]

##-- tmle (section 3.1);
tmle.fit <- tmle(Y=sim.data$Y, A=sim.data$A,
  cbind(X1=sim.data$X1,X2=sim.data$X2,X3=sim.data$X3),
  gform=A~X1+X2+X3, ## treatment model
  Qform=Y~A+X1+X2+X3, ## outcome model
  family="binomial",
  cvQinit=FALSE)
fit.tmle[[m]] <- tmle.fit$estimates$ATE$psi

##-- one-step estimator
sim.data[, pred.f1:=predict(fit.f, type="response", newdata=copy(sim.data)[, A
:=1])]
sim.data[, pred.f0:=predict(fit.f, type="response", newdata=copy(sim.data)[, A
:=0])]
sim.data[, pred.f:=predict(fit.f, type="response", newdata=sim.data)]
sim.data[, pred.pi1:=predict(fit.pi, type="response", newdata=sim.data)]
fit.one.step[[m]] <- sim.data[, mean((A/pred.pi1 - (1-A)/(1-pred.pi1))*(Y - pred.f
) +
  pred.f1 - pred.f0)]
}

```

Task 9. See Figure 1. Note that I have added the one-step estimator to show equivalence to TMLE.

```

setwd("~/Undervisning/TMLE/beamer/day1/")
library(ggplot2)

pdat <- data.table(estimator=c(rep("g-formula estimator (misspecified)",
  length(fit.g.glm1)),
  rep("g-formula estimator (correctly specified)",
  length(fit.g.glm2)),
  rep("g-formula estimator (random forest)",
  length(fit.g.rf)),
  rep("IPW estimator (correctly specified)",
  length(fit.ipw)),
  rep("TMLE estimator (misspecified initial)",
  length(fit.tmle)),
  rep("One-step estimator (misspecified initial)",
  length(fit.one.step))),
  est=c(unlist(fit.g.glm1),
  unlist(fit.g.glm2),
  unlist(fit.g.rf),
  unlist(fit.ipw),
  unlist(fit.tmle),
  unlist(fit.one.step)))

ggplot(pdat) +
  theme_bw(base_size=25) +
  geom_boxplot(aes(x=est)) +
  facet_wrap(. ~ estimator, ncol=2) +

```

```
geom_vline(aes(xintercept=ATE), linetype="dashed", color="red") +  
xlab(expression(hat(psi)[n])) + ylab("")
```

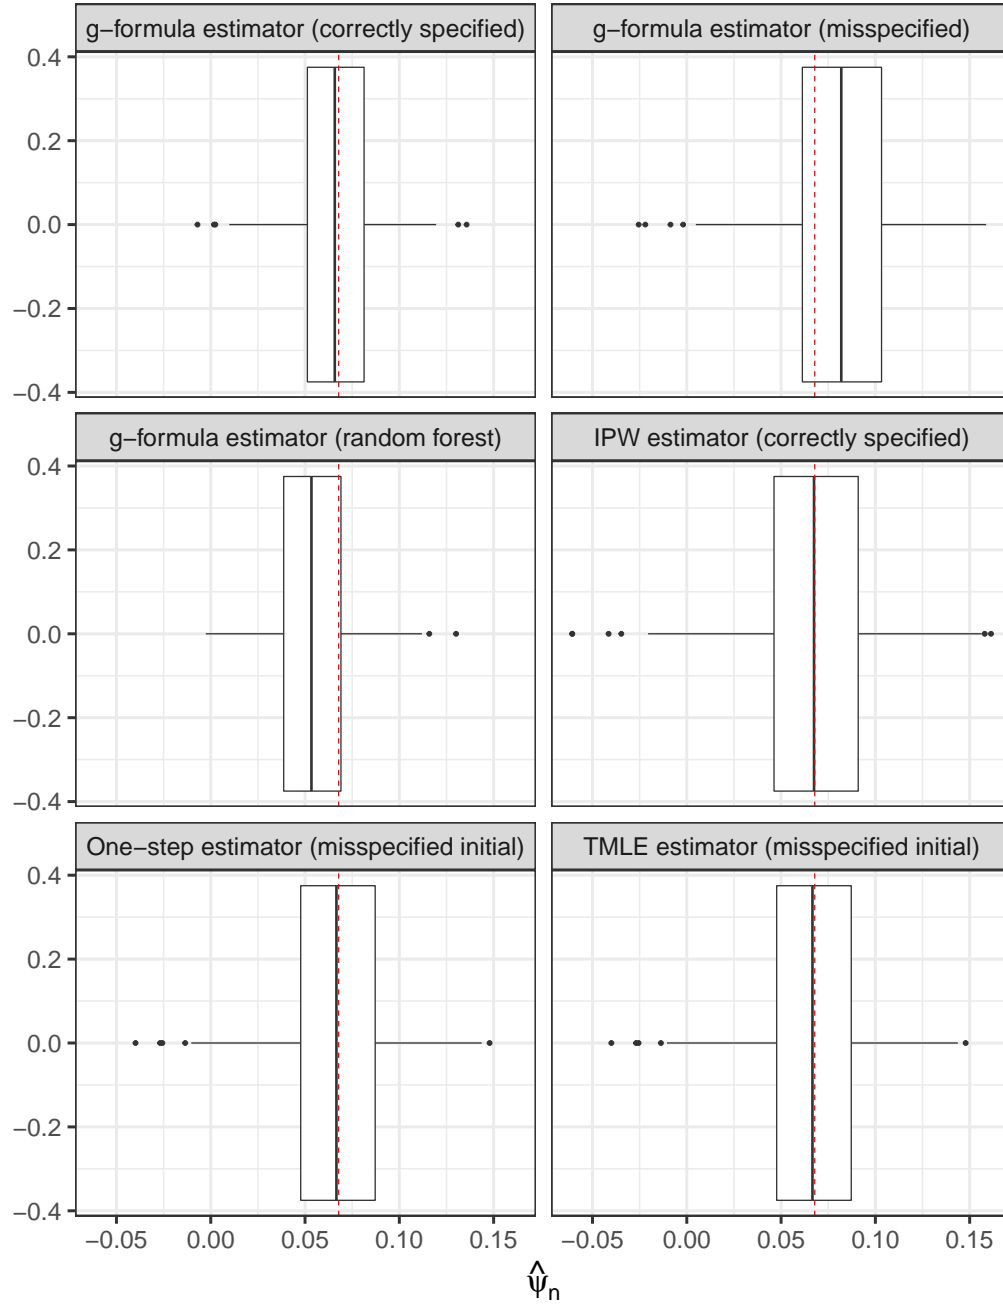


Figure 1

4 Simulation study with positivity issues

Extra for task 9. Same as **Task 9** but changing the distribution of A as follows: $\mathbb{E}[A \mid X_1, X_2, X_3] = \text{logit}(-0.25 + 2.8X_1 + 0.25X_3)$. See Figure 3. Note that I have added the one-step estimator to show better finite-sample distribution of TMLE.

```
new.sim.fun <- function(n, a=NULL) {
  X1 <- runif(n, -2, 2)
  X2 <- rnorm(n)
  X3 <- rbinom(n, 1, 0.2)
  if (length(a)>0) {
    A <- a
  } else {
    A <- rbinom(n, 1, prob=plogis(-0.25 + 2.8*X1 + 0.25*X3))
  }
  Y <- rbinom(n, 1, prob=plogis(-0.9 + 1.9*X1^2 + 0.6*X2 + 0.5*A))
  if (length(a)>0) {
    return(mean(Y))
  } else {
    return(data.table(id=1:n, X1=X1, X2=X2, X3=X3, A=A, Y=Y))
  }
}

fit.g.glm1 <- list()
fit.g.glm2 <- list()
fit.g.rf <- list()
fit.ipw <- list()
fit.tmle <- list()
fit.one.step <- list()

for (m in 1:500) {

  set.seed(m+110)
  sim.data <- new.sim.fun(1000)

  fit.f <- glm(Y~A+X1+X2+X3, family=binomial, data=sim.data)
  fit.f2 <- glm(Y~A+X1.squared+X2+X3, family=binomial, data=sim.data[, X1.squared:=
X1^2])
  fit.rf <- rfsrc(Y~A+X1+X2+X3, data=sim.data)
  fit.pi <- glm(A~X1+X2+X3, family=binomial, data=sim.data)

  ##-- g-formula (section 3.1);
  sim.data[, pred.f1:=predict(fit.f, type="response", newdata=copy(sim.data)[, A
:=1])]
  sim.data[, pred.f0:=predict(fit.f, type="response", newdata=copy(sim.data)[, A
:=0])]
  fit.g.glm1[[m]] <- sim.data[, mean(pred.f1 - pred.f0)]

  ##-- g-formula (section 3.2);
  sim.data[, pred.f1:=predict(fit.f2, type="response", newdata=copy(sim.data)[, A
:=1])]
  sim.data[, pred.f0:=predict(fit.f2, type="response", newdata=copy(sim.data)[, A
:=0])]
  fit.g.glm2[[m]] <- sim.data[, mean(pred.f1 - pred.f0)]
}
```

```

##-- g-formula based on RF (section 3.3);
sim.data[, pred.f1:=predict(fit.rf, type="response", newdata=copy(sim.data)[, A
:=1])$predicted]
sim.data[, pred.f0:=predict(fit.rf, type="response", newdata=copy(sim.data)[, A
:=0])$predicted]
fit.g.rf[[m]] <- sim.data[, mean(pred.f1 - pred.f0)]

##-- ipw (section 3.1);
sim.data[, pred.pi1:=predict(fit.pi, type="response", newdata=sim.data)]
fit.ipw[[m]] <- sim.data[, mean(A*Y/pred.pi1 - (1-A)*Y/(1-pred.pi1))]

##-- tmle (section 3.1);
tmle.fit <- tmle(Y=sim.data$Y, A=sim.data$A,
               cbind(X1=sim.data$X1,X2=sim.data$X2,X3=sim.data$X3),
               gform=A~X1+X2+X3, ## treatment model
               Qform=Y~A+X1+X2+X3, ## outcome model
               family="binomial",
               gbound=0,
               cvQinit=FALSE)
fit.tmle[[m]] <- tmle.fit$estimates$ATE$psi

##-- one-step estimator
sim.data[, pred.f1:=predict(fit.f, type="response", newdata=copy(sim.data)[, A
:=1])]
sim.data[, pred.f0:=predict(fit.f, type="response", newdata=copy(sim.data)[, A
:=0])]
sim.data[, pred.f:=predict(fit.f, type="response", newdata=sim.data)]
sim.data[, pred.pi1:=predict(fit.pi, type="response", newdata=sim.data)]
fit.one.step[[m]] <- sim.data[, mean((A/pred.pi1 - (1-A)/(1-pred.pi1))*(Y - pred.f
) +
                                pred.f1 - pred.f0)]
}

```

```

setwd("~/Undervisning/TMLE/beamer/day1/")
library(ggplot2)

pdat <- data.table(estimator=c(rep("g-formula estimator (misspecified)",
length(fit.g.glm1)),
rep("g-formula estimator (correctly specified)",
length(fit.g.glm2)),
rep("g-formula estimator (random forest)",
length(fit.g.rf)),
rep("IPW estimator (correctly specified)",
length(fit.ipw)),
rep("TMLE estimator (misspecified initial)",
length(fit.tmle)),
rep("One-step estimator (misspecified initial)",
length(fit.one.step))),
est=c(unlist(fit.g.glm1),
unlist(fit.g.glm2),
unlist(fit.g.rf),
unlist(fit.ipw),
unlist(fit.tmle),

```

```

        unlist(fit.one.step)))

ggplot(pdat) +
  theme_bw(base_size=25) +
  geom_boxplot(aes(x=est)) +
  facet_wrap(. ~ estimator, ncol=2) +
  geom_vline(aes(xintercept=ATE), linetype="dashed", color="red") +
  xlab(expression(hat(psi)[n])) + ylab("")

```

```

setwd("~/Undervisning/TMLE/beamer/day1/")
library(ggplot2)

pdat <- data.table(estimator=c(rep("g-formula estimator (misspecified)",
  length(fit.g.glm1)),
  rep("g-formula estimator (correctly specified)",
  length(fit.g.glm2)),
  rep("g-formula estimator (random forest)",
  length(fit.g.rf)),
  rep("IPW estimator (correctly specified)",
  length(fit.ipw)),
  rep("TMLE estimator (misspecified initial)",
  length(fit.tmle)),
  rep("One-step estimator (misspecified initial)",
  length(fit.one.step))),
  est=c(unlist(fit.g.glm1),
  unlist(fit.g.glm2),
  unlist(fit.g.rf),
  unlist(fit.ipw),
  unlist(fit.tmle),
  unlist(fit.one.step)))

ggplot(pdat[estimator %in% c("TMLE estimator (misspecified initial)",
  "One-step estimator (misspecified initial)"]]) +
  theme_bw(base_size=25) +
  geom_boxplot(aes(x=est)) +
  facet_wrap(. ~ estimator, ncol=2) +
  geom_vline(aes(xintercept=ATE), linetype="dashed", color="red") +
  xlab(expression(hat(psi)[n])) + ylab("")

```

```

message(paste0("mse tmle: ", mean((unlist(fit.tmle) - mean(unlist(fit.tmle)))^2)))
message(paste0("mse one step: ", mean((unlist(fit.one.step) - mean(unlist(fit.one.step)
  )))^2)))
message(paste0("variance tmle: ", var(unlist(fit.tmle))))
message(paste0("variance one step: ", var(unlist(fit.one.step))))
message(paste0("bias tmle: ", mean(unlist(fit.tmle)-ATE)))
message(paste0("bias one step: ", mean(unlist(fit.one.step)-ATE)))

```

```

mse tmle: 0.00504335124157197
mse one step: 0.0055914632575187
variance tmle: 0.00505345815788774
variance one step: 0.00560266859470812
bias tmle: 0.00498946280696468
bias one step: 0.00403961703692303

```

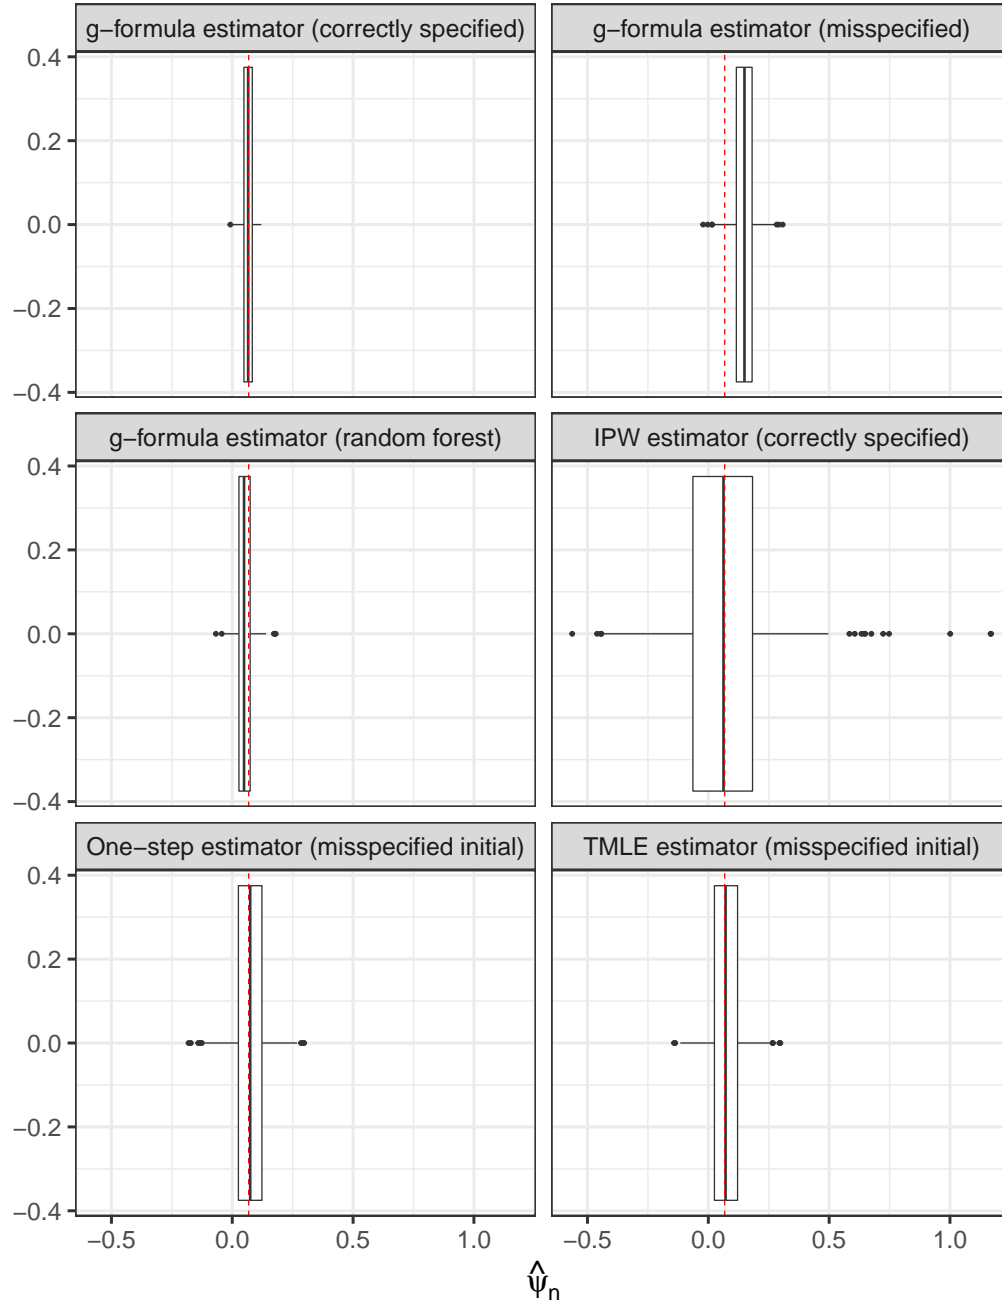


Figure 2

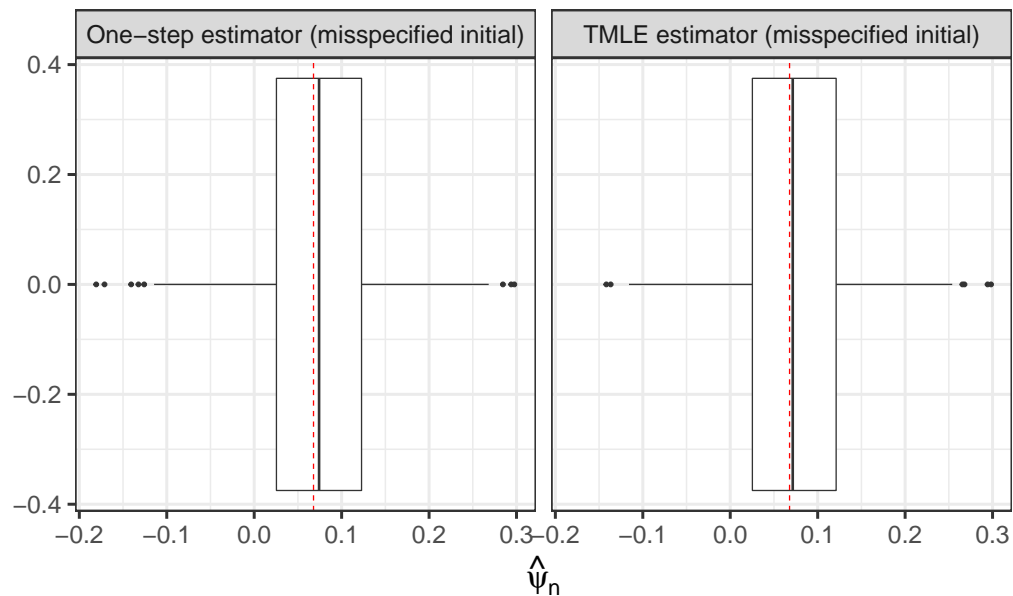


Figure 3

5 Simulation study: Tuning the random forest

```
setwd("~/Undervisning/TMLE/beamer/day1/")
library(ggplot2)
library(data.table)
library(randomForestSRC)

ate.rf.models.list <- list()
V <- 5 # number of folds.
loss.fun <- function(Y, fit) -Y*log(fit)-(1-Y)*log(1-fit) # loss function.

for (m in 1:500) {

  set.seed(5+m)
  sim.data <- sim.fun(1000)
  sim.data[, id:=1:.N]

  #-- what random forests do we want to consider?
  rf.models <- list(rf=c(nodesize=7, mtry=1),
                    rf=c(nodesize=5, mtry=1),
                    rf=c(nodesize=10, mtry=2),
                    rf=c(nodesize=10, mtry=1),
                    rf=c(nodesize=7, mtry=2)
                  )

  #-- for cross-validation;
  cv.split <- matrix(sample(1:nrow(sim.data), size=nrow(sim.data)), ncol=V)

  for (kk in 1:length(rf.models)) {
```

```

rf.model <- rf.models[[kk]]
for (vv in 1:V) {
  test.set <- cv.split[,vv]
  train.set <- sim.data[, id][!sim.data[, id] %in% test.set]
  sim.data.train <- sim.data[id%in%train.set]
  sim.data.test <- sim.data[id%in%test.set]
  train.fit <- rfsrc(formula(paste0("Y~A+X1+X2+X3")),
    data=sim.data.train,
    nodesize=rf.model["nodesize"],
    mtry=rf.model["mtry"])
  sim.data[id%in%test.set, (paste0("fit", kk)):=
    predict(train.fit,
      newdata=sim.data[id%in%test.set],
      type="response")$predicted]
}
}

#-- compute cv error;
cve.rf.models <- unlist(lapply(1:length(rf.models), function(kk) {
  sum(loss.fun(sim.data$Y, sim.data[, get(paste0("fit", kk))]))
})))

#-- fit all random forest models;
for (kk in 1:length(rf.models)) {
  rf.model <- rf.models[[kk]]
  fit.rf <- rfsrc(formula(paste0("Y~A+X1+X2+X3")),
    data=sim.data,
    nodesize=rf.model["nodesize"],
    mtry=rf.model["mtry"])
  sim.data[, (paste0("pred.rf.", kk, ".A1")):=predict(fit.rf, type="response",
    newdata=copy(sim.data)[, A:=1])$predicted]
  sim.data[, (paste0("pred.rf.", kk, ".A0")):=predict(fit.rf, type="response",
    newdata=copy(sim.data)[, A:=0])$predicted]
}

cv.picked <- (1:length(rf.models))[(cve.rf.models==min(na.omit(cve.rf.models))) &
  !is.na(cve.rf.models)]

ate.rf.models.list[[m]] <- c(sapply(1:length(rf.models), function(kk) sim.data[,
  mean(get(paste0("pred.rf.", kk, ".A1"))-
    get(paste0("pred.rf.", kk, ".A0")))]),
  sim.data[, mean(get(paste0("pred.rf.", cv.picked, ".A1"))-
    get(paste0("pred.rf.", cv.picked, ".A0")))]])
}

```




Figure 4